

## 0.1 MISC

### 0.1.1 Alternative QP Formulations

We can take this process one step further and represent the higher-order potential as with non-negativity constraints on  $\theta_k''$  for  $k = 2, \dots, K$ , and appropriate feature vectors, i.e.,

Here we are encoding the coefficients of the pseudo-Boolean function used during inference directly into the learning problem. Like the previous formulation we know that the optimal  $b_1$  is zero so can simply drop  $\theta_0$  and  $\phi_0$  from the optimization.

It is interesting to note the resemblance of the latter QP formulation with latent-variable structural SVM learning [? ].

In our formulation the auxiliary variables  $z$  (see Section 0.2.2) can be determined directly from the ground-truth or inferred labels  $\mathbf{y}$ . Moreover, since we have fixed the piecewise-linear approximation to have equally spaced break-points, the auxiliary variables are independent of the parameters  $(a_k, b_k)$  given  $\mathbf{y}$ . We also have that the  $b_k$  are a function of the  $a_k$  (by Equation ??). Removing the restriction of equally spaced break-points (and introducing the  $b_k$  into the optimization) results in a latent-variable SVM. The main difficulty is that the latent variables  $z$  now depend on the parameters making the optimization problem non-convex.

A number of other variants can be considered by linearly constraining  $\theta$  (or alternatively re-defining  $\phi(\mathbf{y})$ ). For example, the parameters of the  $P^n$ -model can be learned by constraining  $\theta_0 \leq \theta_1$  and forcing  $\theta_i = \theta_{i-1}$  for  $i = 2, \dots, K-1$ . Although this case is somewhat uninteresting as there is only one parameter to learn (since by Observation ?? we can set  $\theta_1 = \dots = \theta_K = 0$  without changing the shape of the potential function), which can often be done more efficiently by other means, e.g., cross-validation over a range of values.

### 0.1.2 Algorithm

**Theorem 0.1.1.:** For the setting  $\epsilon = 0$ , Algorithm ?? terminates with the optimal parameters  $\theta^*$  for  $\text{MAXMARGINQP}(\{\mathbf{y}_t, \mathcal{Y}_t\}_{t=1}^T, \mathbf{D}^2, \mathbf{0})$ .

*Proof.* By Theorem ??, our test for the most violated constraints (lines 7 and 8) can be performed exactly ( $\Delta(\mathbf{y}, \mathbf{y}_t)$  decomposes as a sum of unary terms). If the test succeeds, then  $\mathbf{y}_t^*$  cannot already be in  $\mathcal{A}_t$ . It is now added (line 9). Since there are only finitely many constraints, this happens at most  $2^n - 1$  times (per training example), and the algorithm must eventually terminate. On termination there are no more violated constraints, hence the parameters are optimal.  $\square$

Unfortunately, as our proof suggests, it may take exponential time for the algorithm to reach convergence with  $\epsilon = 0$ . ? ] showed, however, that for  $\epsilon > 0$  and no additional linear constraints (i.e.,  $\mathbf{G} = \mathbf{0}, \mathbf{h} = \mathbf{0}$ ) max-margin learning within a dual optimization framework will terminate in a polynomial number of iterations. Their

result can be extended to the case of additional linear constraints (see the Appendix for details).

### 0.1.3 Max-margin Learning

Given an energy function  $E(\mathbf{y}; \boldsymbol{\theta}) = \boldsymbol{\theta}^T \boldsymbol{\phi}(\mathbf{y})$  parameterized as a linear combination of features  $\boldsymbol{\phi}(\mathbf{y}) \in \mathbb{R}^m$  and weights  $\boldsymbol{\theta} \in \mathbb{R}^m$ , and a set of  $T$  training examples  $\{\mathbf{y}_t\}_{t=1}^T$  the max-margin framework is a principled approach to learning the weights of the model.

In our formulation we will allow additional linear constraints to be imposed on the weights of the form  $\mathbf{G}\boldsymbol{\theta} \geq \mathbf{h}$ , where  $\mathbf{G} \in \mathbb{R}^{d \times m}$  and  $\mathbf{h} \in \mathbb{R}^d$ . This is not typically necessary for max-margin learning, but, as we will see below, is required for enforcing concavity when learning lower linear envelope potentials.

Now, let  $\mathcal{Y}_t = \{0, 1\}^n$  be the set of all possible assignments for the  $t$ -th training example. The (margin-rescaling) max-margin approach formulates learning as a quadratic programming optimization problem,  $\text{MAXMARGINQP}(\{\mathbf{y}_t, \mathcal{Y}_t\}_{t=1}^T, \mathbf{G}, \mathbf{h})$ :

$$\begin{aligned} & \text{minimize} \quad \frac{1}{2} \|\boldsymbol{\theta}\|^2 + \frac{C}{T} \sum_{t=1}^T \xi_t \\ & \text{subject to} \\ & \quad \boldsymbol{\theta}^T \delta \boldsymbol{\phi}_t(\mathbf{y}) + \xi_t \geq \Delta(\mathbf{y}, \mathbf{y}_t), \quad \forall t, \mathbf{y} \in \mathcal{Y}_t, \\ & \quad \xi_t \geq 0, \quad \forall t, \\ & \quad \mathbf{G}\boldsymbol{\theta} \geq \mathbf{h} \end{aligned} \tag{1}$$

where  $\delta \boldsymbol{\phi}_t(\mathbf{y}) \triangleq (\boldsymbol{\phi}_t(\mathbf{y}) - \boldsymbol{\phi}_t(\mathbf{y}_t))$  is the difference between feature representations for some assignment  $\mathbf{y}$  and the  $t$ -th ground-truth assignment  $\mathbf{y}_t$ ,  $C > 0$  is a regularization constant, and  $\Delta(\mathbf{y}, \mathbf{y}_t)$  measures the loss between a ground-truth assignment  $\mathbf{y}_t$  and any other assignment. In our work we use the Hamming loss, which measures the proportion of variables whose corresponding assignments disagree. More formally, the Hamming loss is defined as  $\Delta(\mathbf{y}, \mathbf{y}') = \frac{1}{n} \sum_{i=1}^n \mathbb{I}[y_i \neq y'_i]$ , where  $\mathbb{I}[P]$  is the indicator function taking value one when  $P$  is true and zero otherwise.

The number of constraints in the QP is exponential in the number of variables, and a standard approach to solving the max-margin QP is by adding constraints incrementally. Briefly, at each iteration the algorithm checks for the most violated constraint (for each training example), using *loss-augmented inference*, and, if found, adds it to the constraint set. The algorithm terminates when no more violated constraints are found (see Algorithm ??).

### 0.1.4 Weighted Lower Linear Envelope Potentials

Suppose we have a binary MRFs  $\mathbf{y} = \{y_1, \dots, y_n\}$ ,  $y_i \in \{0, 1\}$ . A higher-order potential  $\psi_c^H(\mathbf{y}_c)$  is an arbitrary function defined on cliques  $\mathbf{y}_c = \{y_i : i \in c\}$  where  $c \subseteq \{1, \dots, n\}$ . Gould[?] proposed it with a weighted lower linear envelope potential expression

$$\psi_c^H(\mathbf{y}_c) \triangleq \min_{k=1,\dots,K} \left\{ a_k W_c(\mathbf{y}_c) + b_k \right\} \quad (2)$$

where  $(a_k, b_k) \in \mathbb{R}^2$  are linear function parameters and

$$W_c(\mathbf{y}_c) = \sum_{i \in c} w_i^c y_i$$

where  $c$  is a clique.  $w_i^c$  is a per-variable non-negative weights for every nodes in each clique and satisfies  $\sum_{i \in c} w_i^c = 1$ .

In the literature they explored many important properties of lower linear envelope. They proved that by maintaining the order of variables  $a_k$  and  $b_k$ , the encoding is ensured free of redundant linear functions (Proposition 3.1[? ]):

$$a_k > a_{k+1} \quad (3)$$

$$b_k < b_{k+1} \quad (4)$$

Another important property is that the inferred assignment  $y^*$  is irrelevant of arbitrarily shifting the piece-wise functions set up or down. Let  $\tilde{\psi}_c^H(y_c) = \min_{k=1,\dots,K} a_k W_c(y_c) + b_k + b$ . We can get:

$$\arg \min_{y_c} \psi_c^H(y_c) = \arg \min_{y_c} \tilde{\psi}_c^H(y_c)$$

Therefore, we can shift arbitrary envelope to zero:

$$b_1 = 0 \quad (5)$$

From equation (3) we know that variable  $b_k$  maintaining increasing order, thus  $b_k > 0$  when  $k > 1$ .

To help with exact inference on the envelope, Gould[?] rewrite it into a quadratic pseudo-Boolean function by introducing  $K - 1$  auxiliary binary variables  $\mathbf{z} = z_1, \dots, z_{K-1}$ :

$$E^c(y_c, \mathbf{z}) = a_1 W_c(y_c) + b_1 + \sum_{k=1}^{K-1} z_k ((a_{k+1} - a_k) W_c(y_c) + b_{k+1} - b_k) \quad (6)$$

It is worth to notice that as long as we maintain the linear constraints on variables  $a_k$  and  $b_k$ , the following constraint is automatically satisfied (Observation 3.6[? ]):

$$z_{k+1} \leq z_k \quad (7)$$

In the literature they also proved the submodularity of equation (5) and proposed

a graph-cuts method to perform the exact inference on it. Our work is based on these results.

## 0.2 Related Work

### 0.2.1 MRF

Learning structural objects from unknown probability distribution is becoming popular in recent years. [?] generalized multiclass SVMs [?] to structural SVMs by extending feature vectors to joint feature vectors which map features extracted jointly over input-output pairs to discrete output. The exact maximum a posteriori (MAP) problem thus becomes an NP-hard problem. They overcome this by generalize the hard margin into "soft" margin and found an upper bound of arbitrary loss functions under this formulation.

Based on the previous research, [?] developed latent SVM by introducing a hidden variable into the joint feature vector. They observed a fact that in real world applications hidden variables are usually intermediate results and are not required as an output. With this fact they followed Soft-Margin method and found an upper bound for the loss function with latent variables. However, the resulted object function is still non-convex.

[?] developed the Concave-Convex Procedure (CCCP) which is guaranteed to find a local minimum for a Difference-Convex (DC) program. [?] combined CCCP algorithm by writing their non-convex object function into a difference of two convex functions and came up with an EM like 2 steps optimizing algorithm. For each iteration, they first compute latent variables utilizing current parameter vectors and then in turn optimizing parameter vectors using the standard Structural SVM algorithm with previously computed latent variables.

Higher order potentials are raising interests due to their capability to represent dependencies between complex objects. [?] proposed a method to represent a class of higher order potentials with lower (upper) linear envelope potentials. By introducing auxiliary variables, they reduced the linear representation to a pairwise form and proposed an approximate algorithm with standard linear programming methods. Following their routine, [?] extended their method to a weighted lower linear envelope in binary Markov Random Fields (MRF) which can be solved with an efficient algorithm. They showed the energy function with auxiliary variables is submodular by transforming it into a quadratic pseudo-Boolean form and how graph-cuts like algorithm can be applied to do exact inference. They then optimized the models parameters under the max margin framework [?].

In their work they pointed out the potential relationship between their auxiliary representation and latent SVM [?]. However, since removing of their fixed space constraint will result dependence between latent variable and parameters, further research still remains open.

This report will study the latent variable formulation of their energy function and optimize it using latent SVM framework. The rest of report is organized as follows:

Section 2 gives a briefly introduction of Latent Structural SVM. Section 3 presents the formulation of weighted lower linear envelope potential and its most important properties. Section 4 describes the linear combination formulation of the energy function and how to optimize it using latent structural SVM. Section 5 discusses its performance on synthetic data.

### 0.2.2 Exact Inference

Importantly,  $E^c(\mathbf{y}_c, \mathbf{z})$  is a *submodular* energy function, which allows us to perform efficient inference by minimizing jointly over both variables  $\mathbf{y}_c$  and auxiliary variables  $\mathbf{z}$ .

It is well known that submodular pairwise energy functions can be minimized exactly in time polynomial in the number of variables by finding the minimum-*st*-cut on a suitably constructed graph [? ? ]. We illustrate one possible construction for  $E^c(\mathbf{y}_c, \mathbf{z})$  in Figure ??.

Using this fact, we can show that an energy function containing arbitrary weighted lower linear envelope potentials can be minimized in polynomial time.

Then  $E(\mathbf{y})$  can be minimized in time polynomial in the number of variables  $n$  and total number of linear envelope functions. Each higher-order term adds  $K - 1$  auxiliary variables so the total number of variables in the augmented energy function is less than  $n$  plus the total number of linear functions.

### 0.2.3 Interactive Figure-Ground Segmentation

We also ran experiments on the real-world “GrabCut” problem introduced by ? ]. Here the aim is to segment a foreground object from an image given a user-annotated bounding box of the object (see for some examples). To solve this problem the GrabCut algorithm associates a binary random variable  $y_i$  with each pixel in the image indicating whether the pixel belongs to the “background” (binary label 0) or the “foreground” (binary label 1). Variables corresponding to pixels outside of the user-annotated bounding box are automatically assigned a label of zero (i.e., background). The assignment for the remaining variables, i.e., those within the bounding box, is inferred.

We compare a model with learned higher-order terms against the baseline GrabCut model by performing leave-one-out cross-validation on a standard set of 50 images from ? ]. Following the approach of ? ], our baseline model contains unary and pairwise terms. The unary terms are defined as the log-likelihood from foreground and background Gaussian mixture models (GMMs) over pixel colour and are image-specific. Briefly, the GMMs are initialized by learning foreground and background models from pixels inside and outside the user-annotated bounding box, respectively. Next, the GMMs are used to relabel pixels (within the bounding box) as either foreground or background by taking the label with highest likelihood according to the current parameter settings. Next the parameters of the foreground and background colour models are re-estimated given the new labeling. This process of parameter estimation and re-labeling is repeated until convergence (or a maximum number of

iterations is reached). The final GMMs are used to construct the unary terms.

The pairwise terms encode smoothness between each pixel and its eight neighbours, and are defined as

$$\psi_{ij}^p(y_i, y_j) = \frac{\lambda}{d_{ij}} \mathbb{I}[y_i \neq y_j] \exp \left\{ -\frac{\|x_i - x_j\|^2}{2\beta} \right\} \quad (8)$$

where  $d_{ij}$  is the distance between pixels  $i$  and  $j$ ,  $x_i$  and  $x_j$  are the RGB colour vectors for pixels  $i$  and  $j$ ,  $\beta$  is the average squared-distance between adjacent RGB colour vectors in the image, and  $\lambda$  determines the strength of the pairwise smoothness term. It is the only free parameter in the baseline model and learned by cross-validation.

To construct the higher-order terms, we adopt a similar superpixel-based approach to [15]. First, we over-segment the image into a few hundred superpixels. Here we use the mean-shift segmentation algorithm of [16] but our method does not depend on this choice. The pixels within each superpixel then define a higher-order term, much like the checkerboard squares in our synthetic experiments. Here, however, the higher-order terms are over different sized cliques and there is no guarantee that they should be labeled homogeneously.

We learn the weights for the unary and pairwise potentials and the parameters for a lower linear envelope potential with  $K = 10$  terms using Algorithm 1. We set  $C = 1000$  and ran for a maximum of 100 iterations, however, for most cross-validation folds, the algorithm converged before the maximum number of iterations was reached. The parameters determined at the last iteration were used for testing. Learning took approximately 3 hours per cross-validation fold with the majority of the time spent generating violated constraints for the 49 training images (each typically containing  $640 \times 480$  pixels).

Some example results are shown in Figure 1. The first row shows that our higher-order terms can capture some fine structure such as the cheetah’s tail but it also segments part of the similarly-appearing rock. In the second example, we are able to correctly segment the person’s legs. The third example shows that we are able to segment the petals at the lower part of the rightmost flower, which the baseline model does not. The final example (fourth row) shows that our model is able to remove background regions that have similar appearance to the foreground. However, we can also make mistakes such as parts of the sculpture’s robe. Quantitatively, our method achieves 91.5% accuracy compared to 90.0% for the strong baseline.